



This would be for each API endpoint. The amount of stages is preference, initially we used 5 but found it to be too cumbersome so we landed on 3.

This is run on a schedule from Azure Data Factory. Nightly, hourly, etc...

This is for a batch dataset however a similar principle could be applied by consuming a kafka endpoint in stage 1 & starting a stream.

